

Assignment Project Exam Help

Introduction

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

29 July 2020

① Organisation

② Distributed algorithms

③ Graphs and spanning trees revisited

④ Ba

⑤ Ec

⑥ Ec

⑦ Echo/size algorithm

⑧ Further algorithms

⑨ Project and practical work

⑩ Readings

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Organisation

Assignment Project Exam Help

- For additional details see the DCO and the Canvas Syllabus
<https://courseoutline.auckland.ac.nz/dco/course/COMPSCI/711/1205>
<https://canvas.auckland.ac.nz/courses/45914>

- <https://eduassistpro.github.io>

- A bigger assignment, called project plus re

- Exam: theory of the distributed algorithm

- Assignments: practical implementations, emulations of specific distributed algorithms on a single computer

Distributed algorithms

Assignment Project Exam Help

- **Computing:** computer systems, networking, computer architectures, computer programming, algorithms



<https://eduassistpro.github.io>

- **Distributed computing:** distributed system communication channels, distributed algorithms

- concurrency of components
- paramount messaging time
- lack of global clock in the async case
- independent failure of components

Overlap between parallel and distributed computing

- One litmus test:

- Parallel computing: **tight coupling** between tasks, that cooperate on shared memory

Distributed computing: **loose coupling** between nodes, that

- <https://eduassistpro.github.io>

- In **classical algorithms**, the problem is **encoded** and given to the (one single) processing element

In **parallel algorithms**, the problem is **encoded** and then its chunks are given to the process

- In **distributed algorithms**, the problem is given by the **network** itself and solved by **coordination protocols**

- More:

https://en.wikipedia.org/wiki/Distributed_computing#Parallel_and_distributed_computing

Typical scenario in distributed computing

- computing nodes have local memories and unique IDs

- nodes are arranged in a network: graph, digraph, ring, ..

- neighbouring nodes can communicate by message passing

- <https://eduassistpro.github.io>

- the network topology (size, diameter, nei
other characteristics (e.g. latencies) are o
individual nodes

- the network may or may not dynamically change

- nodes solve parts of a bigger problem or have individual problems but still need some sort of coordination – which is achieved by distributed algorithms

Recall from graph theory

- Graphs, edges, digraphs, arcs, degrees, connected, complete graphs, size, distance, diameter, radius, paths, weights/costs, shortest paths, spanning trees, ...

<https://eduassistpro.github.io>

- min

- diameter = maximum distance, between

- max min

- radius = minimum maximum distance, for any node to any other node (minimum attained at centers)
- min max min

Typical graphs notations

Graph: (V, E)

Assignment Project Exam Help

- V , set of nodes (vertices), $|V| = n$
- E , set of edges
- $d(u, v)$, shortest path from node u to node v
- $e(u)$, eccentricity of node u
- D , **diameter** = maximum eccentricity, over all nodes
- R , **radius** = minimum eccentricity, from any node
- **centre** (node) = node with minimum eccentricity (not unique)

Geodesics examples

Assignment Project Exam Help

- Three distinct spanning trees (rooted at 1) on the same undirected graph

-

- <https://eduassistpro.github.io>

Add WeChat edu_assist_pr

BFS and DFS spanning trees

- **Reflexive edge**: an edge that loops back to the same node

- Default assumption: unless specifically said so, graphs are **irreflexive** (no reflexive edges)

•

<https://eduassistpro.github.io>

- **BFS spanning tree** characterisations

- Each node at graph hop distance d from the tree

- Frond edges only between nodes at dept **one** (thus linking nodes on **different** branches)

- **DFS spanning tree** characterisation

- Frond edges only between between nodes on the **same** branch (a frond links an ancestor with a descendant)

BFS and DFS spanning trees – examples

- Three distinct spanning trees (rooted at 1) on the same undirected graph

Assignment Project Exam Help

- Which one is BFS or DFS (start from root node 1)?

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

- Left is BFS; Right is DFS (when we first go to node 2)
- Middle is neither, but could be BFS, if we start from 2...
- ... or DFS, if we start from 3 or 4

Rounds and steps

Assignment Project Exam Help

- Nodes work in rounds (macro-steps), which have three

<https://eduassistpro.github.io>

③ Send sub-step: send outgoing message

- Note: some algorithms expect null or empty explicit confirmation that nothing was sent

Timing models

① Synchronous models

- nodes work totally synchronised ~ in lock-step
- easiest to develop, but often unrealistic and less efficient

②

<https://eduassistpro.github.io>

- often unrealistic and sometimes impos

③

Partially synchronous models

- some time bounds or guarantees
 - more realistic, but most difficult
- Quiz: guess which models have been first studied?
Heard of Nasreddin Hodja's lamp? ☺

Synchronous model – equivalent versions

Assignment Project Exam Help

- Synchronous model - version 1

all nodes: process takes 1 time unit

- <https://eduassistpro.github.io>^{nits}

- all nodes: process takes 0 time units

Add WeChat [edu_assist_pr](#)

- The second (and equivalent) version ensures that synchronised models are a particular case of asynchronous models

Asynchronous model

- Asynchronous model

Assignment Project Exam Help

- all nodes: process takes 0 time units

each message (individually): transit time (send receive)

<https://eduassistpro.github.io>

- often, a **FIFO** guarantee, which however timing assumption (see next slide)

Add WeChat edu_assist_pro

- Time complexity (worst-case)** = supremum of all possible normalised async runs
 - NOTE: async time complexity \geq sync time complexity as the sync run is just one of the all possible runs

Asynchronous model

- Asynchronous model with FIFO channels

- the FIFO guarantee: faster messages cannot overtake earlier slower messages sent over the same channel

congestion (pileup) may occur and this should be accounted for

<https://eduassistpro.github.io>

delivered after an additional arbitrary d

[Lynch]

Add WeChat edu_assist_pro

- thus, a sequence of n messages last is delivered
- essentially, a FIFO “channel” may not be a simple channel, but need some middleware (to manage the message queue)
- suggestion to develop robust models, who do not rely on any implicit FIFO assumption [Tel]

Nondeterminism

Assignment Project Exam Help

- In general, both models are **non-deterministic**

<https://eduassistpro.github.io>

- However, all executions must arrive to a **valid** (not necessarily the same, but valid)

Add WeChat [edu_assist_pro](#)

- If always same decisions, then the system is called **confluent**

Starting and stopping options

- Starting options

Assignment Project Exam Help

- Single source or centralised or diffusing: start from a designated **initiator** node

<https://eduassistpro.github.io>

- Termination options

Add WeChat edu_assist_pro

- Easy to notice from outside, but how do the know this?

- Sometimes one node (often the initiator) takes the final decision and can notify the rest (usually omitted phase)
- In general, this can be very challenging and requires sophisticated **control algorithms** for **termination detection**

Echo algorithm

- **Echo** is a fundamental diffusing (single source) algorithm, which is the basis for several others

- combines two phases (imagine a stone thrown into a pond)

<https://eduassistpro.github.io>

- a **convergecast**, "bottom-up" or echo phase, which confirms the termination

Add WeChat [edu_assist_pr](#)
parent-to-child links can be built either a
by additional confirmation messages i
broadcast (not shown in the next slides)

- after receiving echo tokens from all its neighbours, the source node **decides the termination** (and can optionally start a third phase, to inform the rest of this termination)

Echo algorithm

Assignment Project Exam Help

- Echo algorithm is an instance of a large family known as **wave**

- <https://eduassistpro.github.io>
BFS spanning tree – Echo is also known as **SyncBFS**


- In the **async mode**, the broadcast phase of **Ec**
spanning tree, but not necessarily a **BFS spa**

Echo Algorithm - Sync

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 0


Messages = 0

Echo Algorithm - Sync

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 1


Messages = 3

Echo Algorithm - Sync

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 2


Messages = 7

Echo Algorithm - Sync

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 3


Messages = 10

Echo Algorithm - Sync

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = $3 \leq 2D + 1$

Messages = $10 = 2|E|$

Echo programs (Tel)

- Each node has a list, Neigh, indicating its neighbours
- There are two different programs: (1) for the initiator, and (2)

- <https://eduassistpro.github.io>
reduced to this)

- With the exception of the initiator's first state
active only after receiving at least one message
idle (passive) between sending and receiving new messages
- Exercise: try to translate the following pseudocodes into a state machine format

Echo program for initiator (Tel)

Assignment Project Exam Help

```
1 let parent = null
2 let rec
3
4 for
5     send tok
6
7 while rec < | Neigh | do
8     receive tok
9     rec += 1
10
11 decide
```

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Echo program for non-initiators (Tel)

```
1 let parent = null
2 let rec = 0
3
4 receive tok
5 parent = q
6 rec
7
8 for q in Neigh \ parent do
9     send tok to q
10
11 while rec < |Neigh| do // cou
12     receive tok // forward and return
13     rec += 1
14
15 send tok to parent
```

Sync vs. Async

- Using the informal complexity measures appearing in the preceding slides (there are others), we conclude

- Sync**: time complexity = $O(D)$, message complexity = $O(|E|)$

•

<https://eduassistpro.github.io>

- However, the runtime complexity measure (drastically)

- Async**: time complexity = $O(|V|)$, message complexity = $O(|E|)$

- Why?


- For the **time complexity**, we take the **supremum** over all possible **normalised** executions (delivery time in $[0, 1]$)

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 0


Messages = 0

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = ϵ

Messages = 3

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Time Units = 2ε


Messages = 4

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 3ϵ


Messages = 6

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 4ϵ


Messages = 7

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 1


Messages = 7

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 2


Messages = 8

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 3


Messages = 9

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Time Units = 4


Messages = 10

Echo Algorithm - Async

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Total Time Units = 4

Time Units on Broadcast = 1

Time Units on Convergecast = $3 = |V| - 1$

Messages = 10

Echo algorithm revisited

- Like other members of the wave algorithm family, Echo can be adapted to determine:

- The size of the network (number of nodes)

<https://eduassistpro.github.io>

- In general, functions which are associated such as (why?)


- A "leader" (e.g. the node with the highest ID)

- A simple "trick": values can be attached to the tokens!
 - Forward token with null/zero value
 - Return token to parent with subtree value

Echo/size program – associativity and commutativity

Assignment Project Exam Help

<https://eduassistpro.github.io>



Non-deterministic but confluent evaluations – all return 6!

- Left: $(1+2)+3$, $(2+1)+3$, $3+(1+2)$, $3+(2+1)$
- Right: $1+(2+3)$, $1+(3+2)$, $(2+3)+1$, $(3+2)+1$

Echo/size program for initiator

Assignment Project Exam Help

```
1 let parent = null
2 let rec = 0
3 let
4
5 for
6     send (tok, 0)
7
8 while rec < | Neigh | do
9     receive (tok, s) // order irrel
10    rec += 1
11    size += s
12
13 decide size
```

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

Echo/size program for non-initiators

```
1 let parent = null
2 let rec = 0
3 let size = 1
4
5 receive (tok, s)
6 parent = q
7 rec
8
9 for q in Neigh \ parent do
10     send (tok, 0) to q
11
12 while rec < |Neigh| do
13     receive (tok, s)           // fan-out tokens: s=0
14     rec += 1                  // fan-in tokens: s=subtree size
15     size += s
16
17 send (tok, size) to parent // only children really contribute
```

Further algorithms

Assignment Project Exam Help

- Besides some basic fundamental algorithms, we intend to

<https://eduassistpro.github.io>

- Byzantine agreement: "the crown jewel algorithms" – Dijkstra prize 2005, 2001

Add WeChat edu_assist_pro

- The relation between Byzantine algorithms
- all these have practical significance and applications

Project and practical work

Assignment Project Exam Help

- Practical work is necessary to properly understand the concepts

•

- <https://eduassistpro.github.io>

better, emulate distributed systems on a si

- Add WeChat edu_assist_pr

- For uniformity and fairness in marking, we u specifically with C# (or F#)
- The final exam is focused on concepts, does not contain "technological" questions (related to .NET)

Prerequisites (cf. 335)

Assignment Project Exam Help

- Familiarity with C#, at least to the level presented in the C# 5.0 Specifications, Chapter Introduction (pp 1–32):

- <https://eduassistpro.github.io>

- Familiarity with the specific API that we will distributed systems

Add WeChat edu_assist_pro

- Familiarity with searching and reading th
- Familiarity with Linqpad <http://www.linqpad.net/>

How to emulate sync and async distributed systems?

Assignment Project Exam Help

-

- <https://eduassistpro.github.io>

- Multi-process emulation by HTTP/RE
Sinatra)

Add WeChat edu_assist_pr

Readings

Assignment Project Exam Help

- Textbooks (in library – also limited previews on Google books)
 - Lynch, Nancy (1996). Distributed Algorithms. Morgan

<https://eduassistpro.github.io/>
978-0-52179-483-1.

Fekkinck, Wan (2013). Distributed Alg
Approach, MIT Press. (Second Edition)

- Research and overview articles (generally overlapping the textbooks) will be indicated for each topic

Readings 2

- My articles (in collaboration) on bio-inspired models for distributed algorithms may offer additional views
 - network discovery

Assignment Project Exam Help

<https://eduassistpro.github.io>

- image processing (stereo, skeletonisation)
- Byzantine agreement

Add WeChat edu_assist_pro

- formal verification
- hard problems (SAT, TSP, QSAT)
- Pre-print versions published in the CDMTCS research reports

<https://www.cs.auckland.ac.nz/staff/cgi-bin/mjd/secondcgi.pl?date>